

# LLM-Assisted Data Extraction (Human-in-the-Loop)

**Note on data.** This problem set uses **synthetic** (simulated) text snippets provided in the code tutorial. The goal is to practice structured extraction with LLMs, auditing, and evaluation—not to make substantive claims about real events.

**LLM requirement.** Use an LLM that is **free to use**. Recommended: run a local open-source model with `Ollama` (e.g., `llama3.1:8b` or `mistral:7b-instruct`). Do *not* use a paid API.

## Conceptual Questions

Please write three to ten sentence explanations for each of the following questions. **You are only required to answer ONE of the two questions below.**

1. Explain why **schema-constrained** extraction (structured JSON fields with explicit missingness rules) can reduce hallucination relative to free-form summarization. Then explain one limitation: a way the model can still produce systematically wrong extractions even when it outputs “valid” JSON.
2. Human-in-the-loop extraction requires evaluation. Explain why **spot-audits** and **precision/recall**-style evaluation are both needed. In your answer, define one failure mode that would be missed by evaluating only on a small gold set, and one failure mode that would be missed by auditing only a small random sample.

## Applied Exercises

Use the code in the week’s code tutorial and the lecture slides to answer the following questions.

3. **Replace the API call with a free LLM and generate structured extractions.** Start from the provided script (`llm_human.py`) and modify it so it uses a **free LLM**.
  - Keep the `EventExtraction` schema and require the model to output a single JSON object that matches the schema.
  - Run extraction for all documents and save the output table to `outputs/extractions_raw.csv`.
  - Report (briefly) which model you used (e.g., `llama3.1:8b`) and the exact prompt you used.

**Implementation hint (Ollama).** If using Ollama, you may call the local server from Python via HTTP. For example:

```
# pip install requests
import requests, json

payload = {
```

```

    "model": "llama3.1:8b",
    "prompt": "Return ONLY valid JSON matching this schema: ...",
    "stream": False
}
resp = requests.post("http://localhost:11434/api/generate", json=payload).json()
text = resp["response"]
record = json.loads(text) # then validate with Pydantic

```

4. **Uncertainty flags + audit sheet (human-in-the-loop).** Using your extracted dataset:
  - Create at least **four** mechanical review flags (examples: low confidence, missing date, missing country, `geo_precision` is `unknown`, or empty actors list).
  - Create a **single** audit sheet CSV (similar to the tutorial) that includes:
    - (a) the raw text,
    - (b) the extracted fields,
    - (c) the evidence quotes, and
    - (d) blank columns for human corrections and failure-mode tags.
  - Fill out the audit sheet for at least **five** documents and tag a failure mode for any incorrect extraction (e.g., `date_missing`, `location_vague`, `event_type_wrong`, `actor_hallucination`).
  - Report two audit statistics:
    - (a) the share of audited rows marked correct, and
    - (b) the most common failure mode you observed (a small frequency table is fine).
5. **Evaluation + prompt iteration.** Use the small gold set in the tutorial to evaluate event-type classification:
  - Compute and report a classification report (precision/recall/F1) for `event_type`.
  - Create **two** prompt variants (e.g., different missingness instructions; stronger evidence requirements; a shorter event-type taxonomy explanation).
  - Re-run extraction and evaluation for both prompts, and present a **small table** comparing at least:
    - (a) macro-F1,
    - (b) accuracy, and
    - (c) number of items flagged for human review.
  - In 6–10 sentences, defend which prompt you would use in a larger project and why. Your answer must reference both (i) quantitative evaluation and (ii) auditing considerations.
6. **Challenge Question (Optional — if you finish early):** Make the pipeline more robust to long documents or ambiguous text. Choose **ONE** option:
  - (a) **Chunking.** Split each document into 2–3 chunks, run extraction per chunk, and then write a short rule-based aggregation step that outputs one final record (e.g., choose the highest-confidence chunk; union actors; keep the most specific location). Discuss 2–4 sentences on why chunking changes failure modes.
  - (b) **Abstention policy.** Add an explicit “abstain” rule: if the model is unsure, it must set `event_type = other` and add an uncertainty flag. Compare (i) event-type macro-F1 and (ii) the review queue size before vs after abstention. Interpret the trade-off (5–8 sentences).